

Yan Zhou (Terry)

📍 Cambridge, MA 02138 | 📞 +1 (857) 288-9283 | ✉️ terryzhou1211@gmail.com
🌐 <https://tz1211.github.io> | 🌐 <https://github.com/tz1211>

EDUCATION

Harvard University

MS in Data Science, **GPA: 3.92**

Cambridge, United States

Sep 2025 – May 2027 (expected)

- **Selected Coursework:** AI Alignment and Safety, Deep Learning, Machine Learning, MLOps, Computing at Scale, Decision Theory, Statistical Inference

London School of Economics and Political Science (LSE)

BSc in Politics and Data Science, **First Class Honors**

London, United Kingdom

Sep 2022 – Jun 2025

- **Selected Coursework:** Artificial Intelligence, Machine Learning, Multivariate Calculus, Linear Algebra, Algorithm and Data Structure, Probability and Distribution Theory, Mathematical Proof and Analysis, Databases

RESEARCH EXPERIENCE

CRISP Group, Harvard Kempner Institute

Advisor: Prof. Demba Ba

Cambridge, United States

Apr 2026 – Present

- Master's thesis: leveraging weight-space structures in LLMs to build concept detectors robust to distribution shift and adversarial evasion.

ML Foundations Group, Harvard

Advisor: Prof. David Alvarez-Melis

Cambridge, United States

Nov 2025 – Present

- Proposed ADAPT, a framework for amortizing distillation across both model size and post-training variants, producing interpolated families of post-trained LLMs from a single base-model distillation run.
- Our method outperforms boomerang distillation baseline on generation and reasoning tasks while enabling adaptive model-size selection for improved compute-accuracy trade-offs.

AI4LIFE Group, Harvard

Advisors: Prof. Hima Lakkaraju, Dr. Shichang Zhang

Cambridge, United States

Sep 2025 – Present

- Investigating context-dependent behavioral modulation in LLMs. Proposed a novel evaluation framework characterizing user personas across 5 attribute dimensions with corresponding behavioral metrics, enabling contextually-aware safety evaluation beyond static benchmarks.
- Augmenting the framework with mechanistic interpretability to study how models' internal representations of users causally modulate refusal behavior.

Data Science Institute, LSE

Advisor: Prof. Jonathan Cardoso-Silva

London, United Kingdom

Jun 2024 – Jun 2025

- Built an LLM chatbot for answering queries from student and faculty about LSE policy and regulations.
- Led backend development utilizing RAG framework and developed an automated pipeline to crawl, preprocess, and embed LSE regulations into a vector database.

Data Science Institute, LSE

Advisor: Prof. Jonathan Cardoso-Silva

London, United Kingdom

Jun 2023 – Oct 2024

- Built ML framework to identify factors that drive MP rebellions in UK Parliament.
- Utilized speech embeddings to estimate political ideological differences, achieving a 20% increase in F1-score for predicting voting behavior.

PUBLICATIONS

- **Yan Zhou**, Sara Kangaslahti, Jonathan Geuter, Nihal V. Nayak, Marco Fumero, Francesco Locatello, David Alvarez-Melis. (2026). *Thinking at the Right Size: Amortized Distillation Across Post-Trained LLMs*. Findings of the Association for Computational Linguistics: EMNLP 2026.
- **Yan Zhou**, Shichang Zhang, Zidi Xiong, Himabindu Lakkaraju. (2026). *User Persona Subspaces Modulate Refusal Behavior in Language Models*. ICML 2026 Workshop on Mechanistic Interpretability.

TEACHING

Teaching Assistant: Harvard CS2881r AI Alignment and Safety
Prof. Boaz Barak

- Graduate-level course on frontier topics of AI safety.

Cambridge, United States
Aug 2026 – Dec 2026

Teaching Assistant: LSE DS205 Advanced Data Manipulation
Prof. Jonathan Cardoso-Silva

- Undergraduate course on advanced data science and data engineering techniques.
- Delivered a lecture on retrieval-augmented generation (RAG) workflows.
- Graded final student projects, providing detailed technical feedback on implementation and methodology.

London, United Kingdom
Jan 2025 – Jun 2025

AWARDS

Harvard OMPP Summer Research Fellowship

2026

TECHNICAL SKILLS

Languages:
Machine Learning:
ML Engineering:
Data & Visualization:

Python (Advanced), R (Advanced), Typescript (Beginner), SQL, $\text{\LaTeX}$
PyTorch, Transformers, Huggingface, Scikit-learn
vLLM, Finetuning, Knowledge Distillation, Quantization, Agents, RAG, Docker,
Kubernetes, Weights & Biases, Git, CI/CD, GCP, AWS, Slurm
NumPy, Pandas, SciPy, Matplotlib, Plotly, Seaborn, Jupyter