

Yan Zhou (Terry)

📍 Cambridge, MA 02138 | 📞 +1 (857) 288-9283 | ✉ terryzhou1211@gmail.com
🌐 <https://tz1211.github.io> | 🌐 <https://github.com/tz1211>

EDUCATION

Harvard University

MS in Data Science, GPA: 3.92

Cambridge, United States

Sep 2025 – May 2027 (expected)

- **Selected Coursework:** AI Alignment and Safety, Deep Learning, Machine Learning, MLOps, Computing at Scale, Decision Theory, Statistical Inference

London School of Economics and Political Science (LSE)

BSc in Politics and Data Science, First Class Honors

London, United Kingdom

Sep 2022 – Jun 2025

- **Selected Coursework:** Artificial Intelligence, Machine Learning, Multivariate Calculus, Linear Algebra, Algorithm and Data Structure, Probability and Distribution Theory, Mathematical Proof and Analysis, Databases

RESEARCH EXPERIENCE

CRISP Group, Harvard Kempner Institute

Advisor: Prof. Demba Ba

Cambridge, United States

Apr 2026 – Present

- Master's thesis: leveraging weight-space structure in LLMs to build concept detectors robust to distribution shift and adversarial evasion.

ML Foundations Group, Harvard

Advisor: Prof. David Alvarez-Melis

Cambridge, United States

Nov 2025 – Present

- Proposed ADAPT, a framework for amortizing distillation across both model size and post-training variants, producing interpolated families of post-trained LLMs from a single base-model distillation run.
- Our method outperforms boomerang distillation baseline on generation and reasoning tasks while enabling adaptive model-size selection for improved compute-accuracy trade-offs.

AI4LIFE Group, Harvard

Advisors: Prof. Hima Lakkaraju, Dr. Shichang Zhang

Cambridge, United States

Sep 2025 – Present

- Investigating context-dependent behavioral modulation in LLMs. Proposed a novel evaluation framework characterizing user personas across 5 attribute dimensions with corresponding behavioral metrics, enabling contextually-aware safety evaluation beyond static benchmarks.
- Extending the framework with mechanistic interpretability by projecting model activations onto user-attribute directions and applying activation steering to causally validate how internal representations mediate behavioral modulation.

Data Science Institute, LSE

Advisor: Prof. Jonathan Cardoso-Silva

London, United Kingdom

Jun 2024 – Jun 2025

- Built an LLM chatbot for answering queries from student and faculty about LSE policy and regulations.
- Led backend development utilizing RAG framework and developed an automated pipeline to crawl, preprocess, and embed LSE regulations into a vector database.

Data Science Institute, LSE

Advisor: Prof. Jonathan Cardoso-Silva

London, United Kingdom

Jun 2023 – Oct 2024

- Built ML framework to identify factors that drive MP rebellions in UK Parliament.
- Utilized speech embeddings to estimate political ideological differences, achieving a 20% increase in F1-score for predicting voting behavior.

PUBLICATIONS

- **Yan Zhou**, Sara Kangaslahti, Jonathan Geuter, Nihal V. Nayak, Marco Fumero, Francesco Locatello, David Alvarez-Melis. (2026). *Thinking at the Right Size: Amortized Distillation Across Post-Trained LLMs*. Under review.
- **Yan Zhou**, Shichang Zhang, Zidi Xiong, Himabindu Lakkaraju. (2026). *User Persona Subspaces Modulate Refusal Behavior in Language Models*. ICML 2026 Workshop on Mechanistic Interpretability.

TEACHING

Teaching Assistant: Harvard CS2881r AI Alignment and Safety
Prof. Boaz Barak

- Graduate-level course on frontier topics of AI safety.

Cambridge, United States
Aug 2026 – Dec 2026

Teaching Assistant: LSE DS205 Advanced Data Manipulation
Prof. Jonathan Cardoso-Silva

- Undergraduate course on advanced data science and data engineering techniques.
- Delivered a lecture on retrieval-augmented generation (RAG) workflows.
- Graded final student projects, providing detailed technical feedback on implementation and methodology.

London, United Kingdom
Jan 2025 – Jun 2025

AWARDS

Harvard OMPP Summer Research Fellowship

2026

TECHNICAL SKILLS

Languages:
Machine Learning:
ML Engineering:
Data & Visualization:

Python (Advanced), R (Advanced), Typescript (Beginner), SQL, $\text{\LaTeX}$
PyTorch, Transformers, Huggingface, Scikit-learn
vLLM, Finetuning, Knowledge Distillation, Quantization, Agents, RAG, Docker,
Kubernetes, Weights & Biases, Git, CI/CD, GCP, AWS, Slurm
NumPy, Pandas, SciPy, Matplotlib, Plotly, Seaborn, Jupyter